

- When conducting the work of Lab 11, we conducted the test that uses the Central Limit Theorem even though the sample size was “small” (i.e.,  $n < 30$ ). It turns out, that how “far off” the  $t$ -test is can be computed using a first-order Edgeworth approximation for the error. Below, we will do this for the the further observations.

(a) Boos and Hughes-Oliver (2000) note that

$$P(T \leq t) \approx F_Z(t) + \underbrace{\frac{\text{skew}}{\sqrt{n}} \frac{(2t^2 + 1)}{6} f_Z(t)}_{\text{error}},$$

where  $f_Z(\cdot)$  and  $F_Z(\cdot)$  are the Gaussian PDF and CDF and skew is the skewness of the data. What is the potential error in the computation of the  $p$ -value when testing  $H_0 : \mu_X = 0; H_a : \mu_X < 0$  using the zebra finch further data?

```
# part a
zebrafinch.data <- read_csv("zebrafinches.csv")
mu0 <- 0
further.data <- zebrafinch.data$further
n <- length(further.data)

# t.test and t.stat
further.t.test <- t.test(x=further.data, mu = mu0, alternative = "less")
(t.further <- further.t.test$statistic[[1]])

## [1] -7.777991

# potential error calculation
error.num <- skewness(further.data) * (2*t.further^2 + 1) * dnorm(t.further)
error.denom <- 6 * sqrt(n)
(potential.error <- error.num/error.denom)

## [1] -1.226006e-13
```

**Solution:** The potential error in the computation of the  $p$ -value is  $-1.2260063 \times 10^{-13}$  when testing  $H_0 : \mu_X = 0; H_a : \mu_X < 0$  using the zebra finch further data.

- Compute the error for  $t$  statistics from -10 to 10 and plot a line that shows the error across  $t$ . Continue to use the skewness and the sample size for the zebra finch further data.

```
# part b
gg.errors <- rep(NA, length.out = 1000)
gg.tvals <- seq(-10, 10, length.out=1000)

# create data for errors (further data)
for (i in 1:length(gg.tvals)){
  num <- skewness(further.data) * (2*gg.tvals[i]^2 + 1) * dnorm(gg.tvals[i])
  denom <- 6 * sqrt(n)
  gg.errors[i] <- num/denom
}

# plot
errors.plot <- ggplot()+
  geom_line(aes(x= gg.tvals, y = gg.errors))+
  theme_bw()+
```

```
ylab("Potential Error")+
xlab("t")+
ggtitle("Potential Error for t, from -10 to 10")+
geom_vline(aes(xintercept = t.further), color = "red")
```

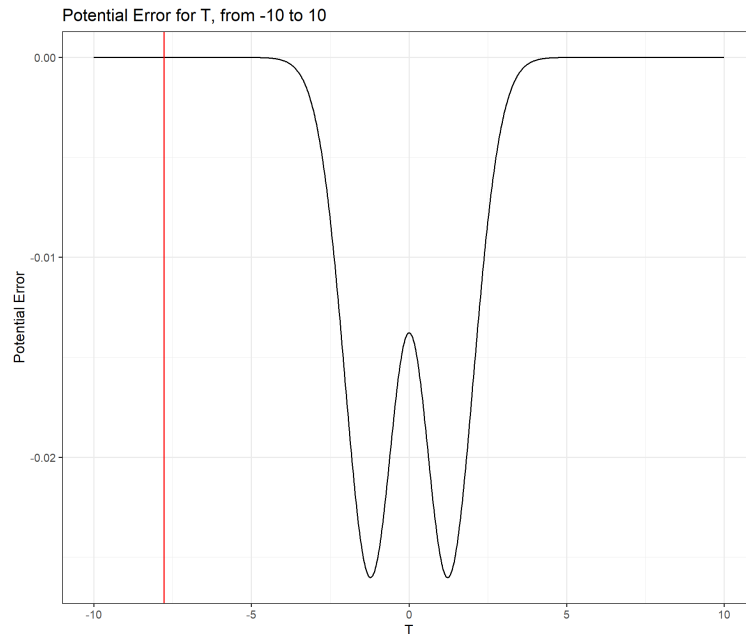


Figure 1: Potential error for  $t$  statistic ranging from -10 to 10, using the skewness of the zebra finch further data. The vertical line shows the  $t$  statistic for the zebra finch further data

**Solution:** Figure 1 plots the error for the  $t$  statistic, ranging from -10 to 10. The  $t$  statistic for the further data is -7.7779912, indicating an exceedingly small error (nearly 0).

- (c) Suppose we wanted to have a tail probability within 10% of the desired  $\alpha = 0.05$ . Recall we did a left-tailed test using the further data. How large of a sample size would we need? That is, we need to solve the error formula equal to 10% of the desired left-tail probability:

$$0.10\alpha \stackrel{\text{set}}{=} \underbrace{\frac{\text{skew}}{\sqrt{n}} \frac{(2t^2 + 1)}{6} f_Z(t)}_{\text{error}},$$

which yields

$$n = \left( \frac{\text{skew}}{6(0.10\alpha)} (2t^2 + 1) f_Z(t) \right)^2.$$

```
alpha <- 0.05
t.alpha <- qnorm(alpha)

(min.nsize <- ((skewness(further.data) * (2*t.alpha^2 + 1) * dnorm(t.alpha))/
  (6 * 0.1 * alpha))^2)

## [1] 520.8876
```

**Solution:** The smallest sample size we would need to have a tail probability within 10% of the desired  $\alpha = 0.05$  is 521. The experiment's  $n = 25$  is significantly smaller than this.

2. Complete the following steps to revisit the analyses from lab 11 using the bootstrap procedure.

- (a) Now, consider the zebra finch data. We do not know the generating distributions for the closer, further, and difference data, so perform resampling to approximate the sampling distribution of the  $T$  statistic:

$$T = \frac{\bar{x}_r - 0}{s/\sqrt{n}},$$

where  $\bar{x}_r$  is the mean computed on the  $r^{\text{th}}$  resample and  $s$  is the sample standard deviation from the original samples. At the end, create an object called `resamples.null.closer`, for example, and store the resamples shifted to ensure they are consistent with the null hypotheses at the average (i.e., here ensure the shifted resamples are 0 on average, corresponding to  $t = 0$ , for each case).

```
# part a
R <- 1000
# data
further.data <- zebrafinch.data$further
closer.data <- zebrafinch.data$closer
diff.data <- zebrafinch.data$diff
# set.seed() to make results reproducible
set.seed(13345)
# resampling
resamples <- tibble(resamps.further = rep(NA, R),
                    resamps.closer = rep(NA, R),
                    resamps.diff = rep(NA, R))
for (i in 1:R){
  # further
  further.resample <- sample(x = further.data,
                            size = length(further.data),
                            replace = T)
  resamples$resamps.further[i] <- (mean(further.resample - mu0))/
    (sd(further.data)/sqrt(n))

  # closer
  closer.resample <- sample(x = closer.data,
                            size = length(closer.data),
                            replace = T)
  resamples$resamps.closer[i] <- (mean(closer.resample - mu0))/
    (sd(closer.data)/sqrt(n))

  # diff
  diff.resample <- sample(x = diff.data,
                          size = length(diff.data),
                          replace = T)
  resamples$resamps.diff[i] <- (mean(diff.resample - mu0))/
    (sd(diff.data)/sqrt(n))
}

# shifting data
resamples <- resamples |>
  mutate(resamps.further.null = resamps.further - mean(resamps.further),
         resamps.closer.null = resamps.closer - mean(resamps.closer),
         resamps.diff.null = resamps.diff - mean(resamps.diff))
```

**Solution:** Above is the code using to conducting resampling for the  $t$  statistic for the zebra finch data. Since we did resampling on the  $t$  statistic, which we want to be centered at 0, we shift the resamples by simply subtracting the mean of all the resamples to center. `set.seed()` was used to create reproducible results.

- (b) Compute the bootstrap  $p$ -value for each test using the shifted resamples. How do these compare to the  $t$ -test  $p$ -values?

```
# part b
boot.pvals <- resamples |>
summarize(p.further = mean(resamps.further.null <= mean(resamps.further)),
          p.closer = mean(resamps.closer.null >= mean(resamps.closer)),
          p.low.diff = mean(resamps.diff.null <= -mean(resamps.diff)),
          p.high.diff = mean(resamps.diff.null >= mean(resamps.diff)),
          pdiff = p.low.diff + p.high.diff)

# p vals from t test
# further
t.pval.further <- t.test(x=further.data, mu = mu0, alternative = "less")$p.value
# closer
t.pval.closer <- t.test(x=closer.data, mu = mu0, alternative = "greater")$p.value
# diff
t.pval.further <- t.test(x=diff.data, mu = mu0, alternative = "two.sided")$p.value

# creating table
all.pvals <- tibble(test = c("T-test", "Bootstrapping"),
                    p.further = c(t.pval.further, boot.pvals$p.further),
                    p.closer = c(t.pval.closer, boot.pvals$p.closer),
                    p.diff = c(t.pval.closer, boot.pvals$pdiff))

xtable(all.pvals)

## % latex table generated in R 4.4.2 by xtable 1.8-4 package
## % Fri May 2 14:24:10 2025
## \begin{table}[ht]
## \centering
## \begin{tabular}{rlrrr}
## \hline
## & test & p.further & p.closer & p.diff \\
## \hline
## 1 & T-test & 0.00 & 0.00 & 0.00 \\
## 2 & Bootstrapping & 0.00 & 0.00 & 0.00 \\
## \hline
## \end{tabular}
## \end{table}
```

	Type	p.further	p.closer	p.diff
1	T-test	0.00	0.00	0.00
2	Bootstrapping	0.00	0.00	0.00

Table 1:  $p$ -values for further, closer, and differences in the zebra finch data using the T-test and Bootstrapping

**Solution:** The  $p$ -values calculated using bootstrapping is so small that it is effectively 0 and treated as such, for the further, closer, and difference zebra finch data. In order to calculate the

two-sided  $p$ -value for the difference data, you need to find the  $p$ -value on both tails, which can be found by adding or subtracting the mean of the difference data by  $\delta$ . However, since  $\mu_0 = 0$ , we can just using the negative of the mean of the difference data to get the “mirror” observation on the other tail.  $p$ -values are compared in Table 1.

- (c) What is the 5<sup>th</sup> percentile of the shifted resamples under the null hypothesis? Note this value approximates  $t_{0.05, n-1}$ . Compare these values in each case.

```
# part c
# 5th percentile for t test (same for all)
t.5th <- qt(p=0.05, df=n-1)

# 5th percentile for bootstrapping
# further
boot.f5th <- quantile(resamples$resamps.further.null, 0.05)

# closer
boot.c5th <- quantile(resamples$resamps.closer.null, 0.05)

# diff
boot.d5th <- quantile(resamples$resamps.diff.null, 0.05)

# create table
all.5thp <- tibble(Type = c("T-test", "Bootstrapping"),
  further = c(t.5th, boot.f5th),
  closer = c(t.5th, boot.c5th),
  diff = c(t.5th, boot.d5th))
xtable(all.5thp)

## % latex table generated in R 4.4.2 by xtable 1.8-4 package
## % Fri May 2 14:24:10 2025
## \begin{table}[ht]
## \centering
## \begin{tabular}{rlrrr}
## \hline
## & Type & further & closer & diff \\
## \hline
## 1 & T-test & -1.71 & -1.71 & -1.71 \\
## 2 & Bootstrapping & -1.60 & -1.58 & -1.59 \\
## \hline
## \end{tabular}
## \end{table}
```

	Type	further	closer	diff
	T-test	-1.71	-1.71	-1.71
	Bootstrapping	-1.60	-1.58	-1.59

Table 2: 5<sup>th</sup> percentile of the shifted resamples of the  $t$  statistic for further, closer, and difference data, compared to the  $t$ -test.

**Solution:** The 5<sup>th</sup> percentile for the  $t$ -test is the same for all three sets of data, as they are all drawn from the  $T$  distribution. The shifted resamples for the  $t$  statistic from bootstrapping was calculated for further, closer, and difference data, and the values are compared in Table 2. Interestingly, all of the bootstrap values are larger than the  $T$ -test ones.

- (d) Compute the bootstrap confidence intervals using the resamples. How do these compare to the  $t$ -test confidence intervals?

```
# part d
# need to keep track of mean (xbar) now
further.xbars <- rep(NA, R)
closer.xbars <- rep(NA, R)
diff.xbars <- rep(NA, R)

for (i in 1:R){
  # further
  further.resample <- sample(x = further.data,
                             size= length(further.data),
                             replace = T)
  further.xbars[i] <- mean(further.resample)

  # closer
  closer.resample <- sample(x = closer.data,
                             size= length(closer.data),
                             replace = T)
  closer.xbars[i] <- mean(closer.resample)

  # diff
  diff.resample <- sample(x = diff.data,
                           size= length(diff.data),
                           replace = T)
  diff.xbars[i] <- mean(diff.resample)
}

# bca test
# helper function
boot.mean <- function(d, i){
  mean(d[i])
}
R <- 1000
# further
further.helper <- boot(data = further.data,
                       statistic = boot.mean,
                       R = R)

# closer
closer.helper <- boot(data = closer.data,
                       statistic = boot.mean,
                       R = R)

# diff
diff.helper <- boot(data = diff.data,
                    statistic = boot.mean,
                    R = R)

# CI
# further
(boot.further.CI <- quantile(further.xbars, c(0.025, 0.975)))

##          2.5%          97.5%
## -0.2540837 -0.1601393
```

```

(t.further.CI <- t.test(x=further.data, mu=mu0, alternative = "two.sided")$conf.int)

## [1] -0.2565176 -0.1489313
## attr("conf.level")
## [1] 0.95

(further.bca <- boot.ci(further.helper, type="bca"))

## BOOTSTRAP CONFIDENCE INTERVAL CALCULATIONS
## Based on 1000 bootstrap replicates
##
## CALL :
## boot.ci(boot.out = further.helper, type = "bca")
##
## Intervals :
## Level      BCa
## 95%      (-0.2596, -0.1578 )
## Calculations and Intervals on Original Scale

# closer
(boot.closer.CI <- quantile(closer.xbars, c(0.025, 0.975)))

##      2.5%      97.5%
## 0.1244849 0.1926684

(t.closer.CI <- t.test(x=closer.data, mu=mu0, alternative = "two.sided")$conf.int)

## [1] 0.1173875 0.1950586
## attr("conf.level")
## [1] 0.95

(closer.bca <- boot.ci(closer.helper, type="bca"))

## BOOTSTRAP CONFIDENCE INTERVAL CALCULATIONS
## Based on 1000 bootstrap replicates
##
## CALL :
## boot.ci(boot.out = closer.helper, type = "bca")
##
## Intervals :
## Level      BCa
## 95%      ( 0.1218,  0.1920 )
## Calculations and Intervals on Original Scale

# diff
(boot.diff.CI <- quantile(diff.xbars, c(0.025, 0.975)))

##      2.5%      97.5%
## 0.2858302 0.4470153

(t.diff.CI <- t.test(x=diff.data, mu=mu0, alternative = "two.sided")$conf.int)

## [1] 0.2719028 0.4459921
## attr("conf.level")
## [1] 0.95

(diff.bca <- boot.ci(diff.helper, type="bca"))

```

```
## BOOTSTRAP CONFIDENCE INTERVAL CALCULATIONS
## Based on 1000 bootstrap replicates
##
## CALL :
## boot.ci(boot.out = diff.helper, type = "bca")
##
## Intervals :
## Level      BCa
## 95%      ( 0.2857,  0.4514 )
## Calculations and Intervals on Original Scale
```

**Solution:** The 95% confidence intervals for the further data are  $(-0.254, -0.160)$  using bootstrap percentile confidence interval,  $(-0.260, -0.158)$  using the bootstrap BCa, and  $(-0.257, -0.149)$  for the t-test.

The 95% confidence intervals for the closer data are  $(0.124, 0.193)$  using bootstrap percentile confidence interval,  $(0.122, 0.192)$  using the bootstrap BCa, and  $(0.117, 0.195)$  for the t-test.

The 95% confidence intervals for the difference data are  $(0.286, 0.447)$  using bootstrap percentile confidence interval,  $(0.286, 0.451)$  using the bootstrap BCa, and  $(0.272, 0.446)$  for the t-test.

Because the CIs were calculated on the data directly, we needed to conduct resampling again for each of the data's sample means instead of t-statistic.

The BCa confidence intervals and the percentile confidence intervals were very similar to each other, suggesting there is not much bias or skewness in the data.

Additionally, since  $n = 25$ , we can assume that at least some sort of normality and symmetry emerges when conducting resampling, as it's close to the boundary for CLT ( $n = 30$ ).

3. Complete the following steps to revisit the analyses from lab 11 using the randomization procedure.
  - (a) Now, consider the zebra finch data. We do not know the generating distributions for the closer, further, and difference data, so perform the randomization procedure.

```
# part a
set.seed(13345)
R <- 1000
rand <- tibble(xbars.further = rep(NA, R),
               xbars.closer = rep(NA, R),
               xbars.diff = rep(NA, R))
# since mean 0 (mu0 = 0) under H0 is given, no need to shift

# RANDOMIZE / SHUFFLE
for(i in 1:R){
  # further
  further.rand <- further.data *
    sample(x = c(-1, 1),
           size = length(further.data),
           replace = T)

  rand$xbars.further[i] <- mean(further.rand)

  # closer
  closer.rand <- closer.data *
    sample(x = c(-1, 1),
           size = length(closer.data),
           replace = T)
```



```

rand$xbars.closer[i] <- mean(closer.rand)

# diff
diff.rand <- diff.data *
  sample(x = c(-1, 1),
        size = length(diff.data),
        replace = T)

rand$xbars.diff[i] <- mean(diff.rand)
}

# no need to shift back either!

```

**Solution:** The procedure for the randomization test is conducted in the code above. However, because our  $H_0$  is specified as  $\mu = 0$ , we do not need to shift the data before and after shuffling because its already centered around 0.

- (b) Compute the randomization test  $p$ -value for each test.

```

# task b
# p-value
(rand.pvals <- rand |>
  summarize(p.further = mean(xbars.further <= mean(further.data)),
            p.closer = mean(xbars.closer >= mean(closer.data)),
            p.low.diff = mean(xbars.diff <= -mean(diff.data)),
            p.high.diff = mean(xbars.diff >= mean(diff.data)),
            pdiff = p.low.diff + p.high.diff))

## # A tibble: 1 x 5
##   p.further p.closer p.low.diff p.high.diff pdiff
##   <dbl>    <dbl>    <dbl>    <dbl> <dbl>
## 1         0         0         0         0     0

```

	Type	p.further	p.closer	p.diff
	T-test	0.00	0.00	0.00
	Bootstrapping	0.00	0.00	0.00
	Randomization	0.00	0.00	0.00

Table 3:  $p$ -values for further, closer, and differences in the zebra finch data using the T-test, Bootstrapping, and Randomization

**Solution:** Similar to the bootstrap  $p$ -values, the values are so small that it is effectively 0 across all the data. The randomization  $p$ -value procedure is nearly identical to the bootstrap  $p$ -value, except we are now comparing shifted resampled means to the sample mean.  $p$ -values are compared in Table 3.

- (c) Compute the randomization confidence interval by iterating over values of  $\mu_0$ .

**Hint:** You can “search” for the lower bound from  $Q_1$  and subtracting by 0.0001, and the upper bound using  $Q_3$  and increasing by 0.0001. You will continue until you find the first value for which the two-sided  $p$ -value is greater than or equal to 0.05.

```

# task c
# further
R <- 1000

```

```

mu0.iterate.further <- 0.01
starting.point.further <- mean(further.data)

mu.lower.further <- starting.point.further
repeat{
  rand <- tibble(xbars = rep(NA, R))

  # PREPROCESSING: shift the data to be mean 0 under H0
  x.shift <- further.data - mu.lower.further
  # RANDOMIZE / SHUFFLE
  for(i in 1:R){
    curr.rand <- x.shift *
      sample(x = c(-1, 1),
             size = length(x.shift),
             replace = T)

    rand$xbars[i] <- mean(curr.rand)
  }
  # Thinking is hard
  rand <- rand |>
    mutate(xbars = xbars + mu.lower.further) # shifting back

  # p-value
  p.val <- mean(rand$xbars >= mean(further.data))

  if(p.val < 0.05){
    break
  }else{
    mu.lower.further <- mu.lower.further - mu0.iterate.further
  }
}

mu.upper.further <- starting.point.further
repeat{
  rand <- tibble(xbars = rep(NA, R))

  # PREPROCESSING: shift the data to be mean 0 under H0
  x.shift <- further.data - mu.upper.further
  # RANDOMIZE / SHUFFLE
  for(i in 1:R){
    curr.rand <- x.shift *
      sample(x = c(-1, 1),
             size = length(x.shift),
             replace = T)

    rand$xbars[i] <- mean(curr.rand)
  }
  # Thinking is hard
  rand <- rand |>
    mutate(xbars = xbars + mu.upper.further) # shifting back

  # p-value
  p.val <- mean(rand$xbars <= mean(further.data))

```

```

if(p.val < 0.05){
  break
}else{
  mu.upper.further <- mu.upper.further + mu0.iterate.further
}
}

c(mu.lower.further, mu.upper.further)

## [1] -0.2527244 -0.1527244

# closer
R <- 1000
mu0.iterate.closer <- 0.01
starting.point.closer <- mean(closer.data)

mu.lower.closer <- starting.point.closer
repeat{
  rand <- tibble(xbars = rep(NA, R))

  # PREPROCESSING: shift the data to be mean 0 under H0
  x.shift <- closer.data - mu.lower.closer
  # RANDOMIZE / SHUFFLE
  for(i in 1:R){
    curr.rand <- x.shift *
      sample(x = c(-1, 1),
            size = length(x.shift),
            replace = T)

    rand$xbars[i] <- mean(curr.rand)
  }
  # Thinking is hard
  rand <- rand |>
    mutate(xbars = xbars + mu.lower.closer) # shifting back

  # p-value
  p.val <- mean(rand$xbars >= mean(closer.data))

  if(p.val < 0.05){
    break
  }else{
    mu.lower.closer <- mu.lower.closer - mu0.iterate.closer
  }
}

mu.upper.closer <- starting.point.closer
repeat{
  rand <- tibble(xbars = rep(NA, R))

  # PREPROCESSING: shift the data to be mean 0 under H0
  x.shift <- closer.data - mu.upper.closer
  # RANDOMIZE / SHUFFLE
  for(i in 1:R){
    curr.rand <- x.shift *

```

```

    sample(x = c(-1, 1),
           size = length(x.shift),
           replace = T)

    rand$xbars[i] <- mean(curr.rand)
  }
  # Thinking is hard
  rand <- rand |>
    mutate(xbars = xbars + mu.upper.closer) # shifting back

# p-value
p.val <- mean(rand$xbars <= mean(closer.data))

if(p.val < 0.05){
  break
}else{
  mu.upper.closer <- mu.upper.closer + mu0.iterate.closer
}
}

c(mu.lower.closer, mu.upper.closer)

## [1] 0.1162231 0.1962231

# diff
R <- 1000
mu0.iterate.diff <- 0.01
starting.point.diff <- mean(diff.data)

mu.lower.diff <- starting.point.diff
repeat{
  rand <- tibble(xbars = rep(NA, R))

  # PREPROCESSING: shift the data to be mean 0 under H0
  x.shift <- diff.data - mu.lower.diff
  # RANDOMIZE / SHUFFLE
  for(i in 1:R){
    curr.rand <- x.shift *
      sample(x = c(-1, 1),
             size = length(x.shift),
             replace = T)

    rand$xbars[i] <- mean(curr.rand)
  }
  # Thinking is hard
  rand <- rand |>
    mutate(xbars = xbars + mu.lower.diff) # shifting back

# p-value
delta <- abs(mean(diff.data) - mu.lower.diff)
low <- mu.lower.diff - delta # mirror
high <- mu.lower.diff + delta # xbar
p.val <- mean(rand$xbars <= low) +
  mean(rand$xbars >= high)

```

```

if(p.val < 0.05){
  break
}else{
  mu.lower.diff <- mu.lower.diff - mu0.iterate.diff
}
}

mu.upper.diff <- starting.point.diff
repeat{
  rand <- tibble(xbars = rep(NA, R))

  # PREPROCESSING: shift the data to be mean 0 under H0
  x.shift <- diff.data - mu.upper.diff
  # RANDOMIZE / SHUFFLE
  for(i in 1:R){
    curr.rand <- x.shift *
      sample(x = c(-1, 1),
             size = length(x.shift),
             replace = T)

    rand$xbars[i] <- mean(curr.rand)
  }
  # Thinking is hard
  rand <- rand |>
    mutate(xbars = xbars + mu.upper.diff) # shifting back

  # p-value
  delta <- abs(mean(diff.data) - mu.upper.diff)
  low <- mu.upper.diff - delta # mirror
  high <- mu.upper.diff + delta # xbar
  p.val <- mean(rand$xbars <= low) +
    mean(rand$xbars >= high)

  if(p.val < 0.05){
    break
  }else{
    mu.upper.diff <- mu.upper.diff + mu0.iterate.diff
  }
}

c(mu.lower.diff, mu.upper.diff)

## [1] 0.2689475 0.4489475

```

**Solution:** The 95% randomization confidence intervals for the further data is  $(-0.253, -0.153)$ , compared to the t-test  $(-0.257, -0.149)$ , the bootstrap BCa  $(-0.260, -0.158)$ , and the bootstrap percentile  $(-0.254, -0.160)$ .

The 95% randomization confidence intervals for the closer data is  $(0.116, 0.196)$ , compared to the t-test  $(0.117, 0.195)$ , the bootstrap BCa  $(0.122, 0.192)$ , and the bootstrap percentile  $(0.124, 0.193)$ . The 95% randomization confidence intervals for the difference data is  $(0.269, 0.449)$ , compared to the t-test  $(0.272, 0.446)$ , the bootstrap Bca  $(0.286, 0.451)$ , and the bootstrap percentile  $(0.286, 0.447)$ . As we can see, the confidence intervals for all the data using randomization is similar to the confidence intervals from the t-test and bootstrapping, which should make sense. `xbars` needs to be shifted when calculating the confidence interval because `mu.lower` and `mu.upper` change as

we loop through to check. The randomization confidence interval is a lot more “brute force” compared to the others, and we have to search over a grid manually to find the them. While we have historically used `R <- 10000` in randomization, I switched to `R <- 1000` so the code runs quicker.

4. **Optional Challenge:** In this lab, you performed resampling to approximate the sampling distribution of the  $T$  statistic using

$$T = \frac{\bar{x}_r - 0}{s/\sqrt{n}}.$$

I’m curious whether it is better/worse/similar if we computed the statistics using the sample standard deviation of the resamples ( $s_r$ ), instead of the original sample ( $s$ )

$$T = \frac{\bar{x}_r - 0}{s_r/\sqrt{n}}.$$

- (a) Perform a simulation study to evaluate the Type I error for conducting this hypothesis test both ways.
- (b) Using the same test case(s) as part (a), compute bootstrap confidence intervals and assess their coverage – how often do we ‘capture’ the parameter of interest?

## References

Boos, D. D. and Hughes-Oliver, J. M. (2000). How large does  $n$  have to be for  $z$  and  $t$  intervals? *The American Statistician*, 54(2):121–128.